

Experiment

This case study serves as an experiment in determining how approximate computation techniques may change the performance of a multilayer perceptron (MLP) neural network classifier. We have chosen to a subsection of the MNIST data set containing 28 x 28 pixel images of handwritten digits, each labeled according the digit that appears in the image. Each pixel is given by an integer 0 through 255 which indicates the gray-scale value. In the case of each image sample, the digit was known to be roughly centered in the figure. Examples of a few images can be found below in fig. (1) . From the full data set, the first 10,000 images were chosen to train each model, and the next 6,000 were chosen to test the model.

Given that each image is centered in a given sample, we propose that approximating a border of pixels around the outside of the image would produce small changes in the performance of the classifier model, given that dominant features would be preserved. To apply this approximation, we tested a baseline set of samples against samples that have varying depth pixel borders that have been approximated. This effectively creates a padding 3, 5, or 7 pixels on each image. Each approximated pixel value was transformed from a byte to a double precision floating-point number, where the most-significant-bit in the float's exponent was forced to 0, thus only allowing those pixels to exist in the closed interval $[-1, +1]$.

For some samples, it can be seen that even approximating a large border of pixels, the characteristics of the digit are preserved, see fig.1(h). For other samples, characteristic information can be lost and lead to classification errors, see fig. (1g). Each training and testing set of digit images was given to different MLP architecture variants. We adjusted the hidden layers to contain 20, 40, 60, 80, 100, and 120 neurons per layer, for one and two hidden layers. 10 training epochs were used in all cases, with batch sizes of 100 samples used in each step for a Stochastic Gradient Descent Optimizer.


Conclusion

The preliminary results of this experiment, for one and two hidden layers, over the selected set of neuron densities shows that applying a 3, 5, and 7 border of approximated pixels presents no major deviation in performance from the baseline models. In fig. (2a) and (2b), the consistency between training and test samples shows no discernible changes in the loss function. This means that while a subset of pixels were approximated the dominant features preserved in the center of the figure still allowed the optimization process to converge on a similar set of parameters as the baseline samples.

Similarly, the classification of previously unseen samples, as judged by each model's classification precision (2c. , 2d.) and recall (2e. , 2f.) scores also differs very slightly from the


baseline. Due the relatively small volume of training training data and such few epochs, these score ranks low across all models, and particularly low as the network depth increases from a single hidden layer to two hidden layers. However the small change in performance from the baseline indicates that charicteristic features are preserved, and the overall performance of the model remains consistent, even with a border of approximated pixels.

Original: 4




(a)

Original: 9




(b)

3 Pixel Border: 4




(c)

3 Pixel Border: 9




(d)

5 Pixel Border: 4




(e)

5 Pixel Border: 9




(f)

7 Pixel Border: 4




(g)


7 Pixel Border: 9



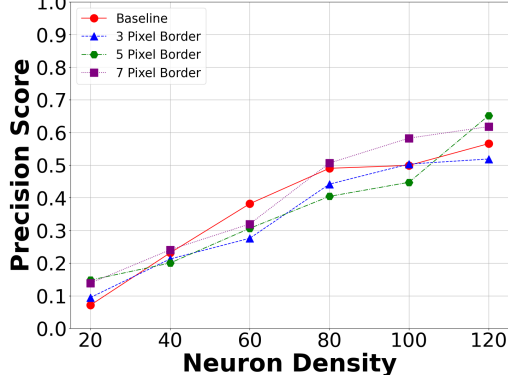
(h)

Single Hidden Layer Loss Value

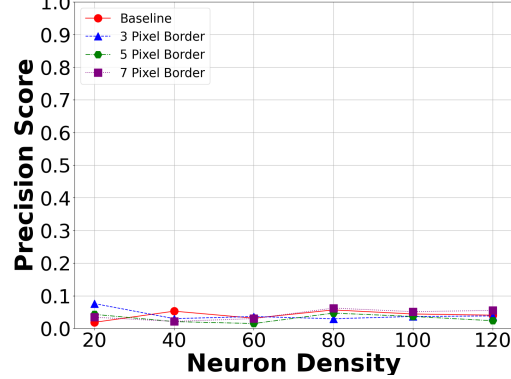
(a)

Double Hidden Layer Loss Value

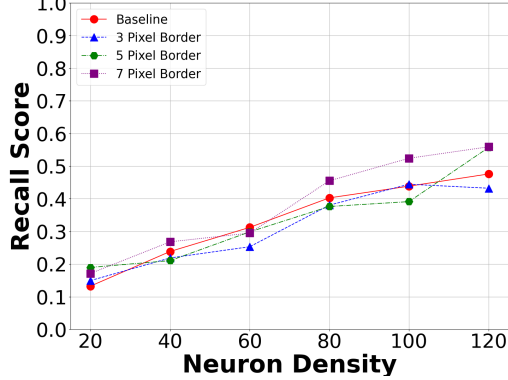
(b)

Single Hidden Layer Precision

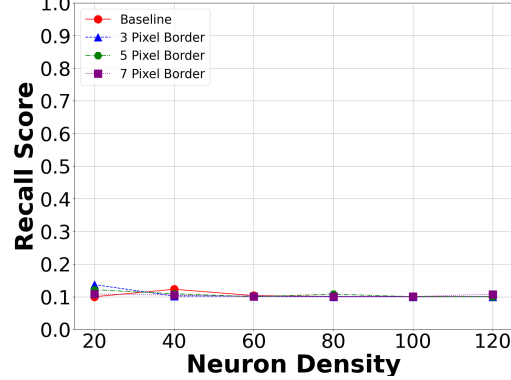
(c)

Double Hidden Layer Precision

(d)

Single Hidden Layer Recall

(e)

Double Hidden Layer Recall

(f)

Figure 2: Experimental results for baseline and 3, 5, and 7 approximated pixel padding border. The same approximations were applied to all samples in both training and testing sets. All data points are taken from the average of 10 models with identical hyper parameters, and non-identical initial weight parameters.